

DT503_3_Traffic_Accident_Data_Analysis_Walkthrough

November 12, 2024

1 Traffic Accident Data Analysis

This notebook explores and analyses a subset of the 'dft-road-casualty-statistics-collision-2023' traffic accident dataset, filtered to a specific region. Using techniques such as Principal Component Analysis (PCA), Factor Analysis, Cluster Analysis and Discriminant Analysis. The goal is to uncover patterns, reduce data dimensionality and develop a reliable classification model.

1.1 Analysis Techniques Covered:

- **Principal Component Analysis (PCA):** for dimensionality reduction
- **Factor Analysis:** to identify latent factors
- **Cluster Analysis:** to find natural groupings in the data
- **Discriminant Analysis:** to develop a classification model

1.2 Step 1: Load and Inspect Data

Let's begin by loading the dataset and examining its structure.

```
[2]: import pandas as pd

# Load the dataset
data = pd.
    ↪read_csv('dft-road-casualty-statistics-collision-2023_Devon_and_Cornwall.
    ↪csv')

# Display the first few rows and basic info
data.info()
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2806 entries, 0 to 2805
```

```
Data columns (total 37 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	accident_index	2806 non-null	float64
1	accident_year	2806 non-null	int64
2	accident_reference	2806 non-null	int64
3	location_easting_osgr	2806 non-null	int64
4	location_northing_osgr	2806 non-null	int64

```

5  longitude                2806 non-null    float64
6  latitude                 2806 non-null    float64
7  police_force             2806 non-null    int64
8  accident_severity        2806 non-null    int64
9  number_of_vehicles        2806 non-null    int64
10 number_of_casualties      2806 non-null    int64
11 date                    2806 non-null    object
12 day_of_week              2806 non-null    int64
13 time                    2806 non-null    object
14 local_authority_district 2806 non-null    int64
15 local_authority_ons_district 2806 non-null    object
16 local_authority_highway  2806 non-null    object
17 first_road_class         2806 non-null    int64
18 first_road_number        2806 non-null    int64
19 road_type                2806 non-null    int64
20 speed_limit              2806 non-null    int64
21 junction_detail          2806 non-null    int64
22 junction_control         2806 non-null    int64
23 second_road_class        2806 non-null    int64
24 second_road_number       2806 non-null    int64
25 pedestrian_crossing_human_control 2806 non-null    int64
26 pedestrian_crossing_physical_facilities 2806 non-null    int64
27 light_conditions         2806 non-null    int64
28 weather_conditions       2806 non-null    int64
29 road_surface_conditions  2806 non-null    int64
30 special_conditions_at_site 2806 non-null    int64
31 carriageway_hazards      2806 non-null    int64
32 urban_or_rural_area      2806 non-null    int64
33 did_police_officer_attend_scene_of_accident 2806 non-null    int64
34 trunk_road_flag          2806 non-null    int64
35 lsoa_of_accident_location 2806 non-null    object
36 enhanced_severity_collision 2806 non-null    int64
dtypes: float64(3), int64(29), object(5)
memory usage: 811.2+ KB

```

```

[2]:  accident_index  accident_year  accident_reference  location_easting_osgr  \
0      2.023500e+12      2023      501258810      251899
1      2.023500e+12      2023      501259203      285979
2      2.023500e+12      2023      501259228      202261
3      2.023500e+12      2023      501259501      251599
4      2.023500e+12      2023      501259642      244791

      location_northing_osgr  longitude  latitude  police_force  \
0          51705 -4.082892  50.346591      50
1          64004 -3.607870  50.464747      50
2          52632 -4.780271  50.340338      50
3          59514 -4.090189  50.416697      50

```

```

4          57992  -4.185314  50.401258          50

    accident_severity  number_of_vehicles  ...  light_conditions  \
0          2          2  ...          1
1          3          2  ...          1
2          3          2  ...          1
3          3          2  ...          4
4          3          2  ...          1

    weather_conditions  road_surface_conditions  special_conditions_at_site  \
0          1          2          0
1          4          2          7
2          5          2          0
3          5          2          0
4          2          2          0

    carriageway_hazards  urban_or_rural_area  \
0          0          2
1          0          2
2          0          1
3          0          1
4          0          1

    did_police_officer_attend_scene_of_accident  trunk_road_flag  \
0          1          2
1          1          2
2          2          2
3          3          2
4          1          2

    lsoa_of_accident_location  enhanced_severity_collision
0          E01015129          6
1          E01020170          3
2          E01019041          3
3          E01015090          3
4          E01015143          3

[5 rows x 37 columns]

```

1.3 Step 2: Principal Component Analysis (PCA)

PCA reduces the data's dimensionality by finding components that explain most of the variance. This helps simplify complex datasets.

```
[3]: from sklearn.decomposition import PCA
     from sklearn.preprocessing import StandardScaler
```

```

# Selecting relevant columns for PCA
numerical_columns = [
    'location_easting_osgr', 'location_northing_osgr', 'accident_severity',
    'number_of_vehicles', 'number_of_casualties', 'day_of_week',
    'first_road_class', 'road_type', 'speed_limit', 'junction_detail',
    'junction_control', 'pedestrian_crossing_human_control',
    'pedestrian_crossing_physical_facilities', 'light_conditions',
    'weather_conditions', 'road_surface_conditions',
    ↪ 'special_conditions_at_site',
    'carriageway_hazards', 'urban_or_rural_area',
    'did_police_officer_attend_scene_of_accident', 'trunk_road_flag',
    'enhanced_severity_collision'
]

# Standardising the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data[numerical_columns])

# Applying PCA
pca = PCA()
pca.fit(data_scaled)
explained_variance_ratio = pca.explained_variance_ratio_
explained_variance_ratio

```

```

[3]: array([0.12133061, 0.10018245, 0.08198755, 0.07498831, 0.07042889,
           0.0641444 , 0.06038357, 0.04953602, 0.04505504, 0.04405497,
           0.04015928, 0.03555746, 0.0319206 , 0.02973638, 0.02878374,
           0.02691546, 0.02580875, 0.01677796, 0.01536993, 0.01344648,
           0.01241099, 0.01102115])

```

1.3.1 PCA Explained Variance


We can visualise the cumulative variance explained by each component.

```

[4]: import matplotlib.pyplot as plt

# Plotting cumulative explained variance
plt.plot(range(1, len(explained_variance_ratio) + 1), pca.
    ↪ explained_variance_ratio_.cumsum(), marker='o')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('PCA - Cumulative Explained Variance')
plt.grid()
plt.show()

```



The cumulative explained variance plot shows the proportion of variance captured by each component.

1.3.2 Interpretation

The original dataset contained 22 numerical components relevant to accident conditions (after selecting and standardising them).

PCA revealed that the first 5 components explained approximately 45% of the total variance in the data. This means that instead of using all 22 original variables, we can focus on these 5 main components for a simpler yet representative view of the dataset.

The reduction from 22 to 5 components enables us to capture the primary patterns without the full complexity of the original data, facilitating clearer analysis in the subsequent steps.

1.4 Step 3: Factor Analysis

Factor Analysis identifies latent factors within the data. These factors help explain the underlying structure of relationships among variables.

1.4.1 Factor Loadings Heatmap

The heatmap below shows the factor loadings, highlighting which variables contribute most to each factor.


```
[5]: from sklearn.decomposition import FactorAnalysis

# Applying Factor Analysis
factor_analysis = FactorAnalysis(n_components=5, random_state=0)
factors = factor_analysis.fit_transform(data_scaled)

# Analyzing factor loadings
factor_loadings = factor_analysis.components_

[6]: import seaborn as sns

# Creating a heatmap of factor loadings
plt.figure(figsize=(12, 6))
sns.heatmap(factor_loadings, annot=True, cmap='coolwarm', yticklabels=[f'Factor_{i+1}' for i in range(factor_loadings.shape[0])],
            xticklabels=numerical_columns)
plt.title('Factor Loadings Heatmap')
plt.show()
```



The heatmap of factor loadings shows the strength of association between each variable and the latent factors identified.

- Each factor (e.g., Factor 1, Factor 2) has strong loadings for certain variables, suggesting specific accident patterns.
- One factor may have high loadings on variables related to road type and junction details, suggesting an underlying pattern linked to road infrastructure.

1.4.2 Interpretation

Factor analysis helps us to identify groups of related variables, enabling a deeper understanding of accident characteristics. Latent factors simplify the dataset by grouping variables with shared patterns, enhancing data interpretation.

1.5 Step 4: Cluster Analysis

Using the factors from the Factor Analysis step, we can apply clustering techniques to segment the data into distinct groups, which may highlight different types of accident profiles.

1.5.1 Cluster Scatter Plot


The scatter plot below shows the clusters on the first two principal components to help visualise how the data is grouped.

```
[7]: from sklearn.cluster import KMeans

# Performing KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=0)
clusters = kmeans.fit_predict(factors)

# Adding the cluster labels to the original data
data['Cluster'] = clusters

# Scatter plot of clusters
plt.figure(figsize=(10, 6))
plt.scatter(factors[:, 0], factors[:, 1], c=clusters, cmap='viridis',
            ↪marker='o', alpha=0.6)
plt.xlabel('Factor 1')
plt.ylabel('Factor 2')
plt.title('Cluster Analysis - Scatter Plot')
plt.colorbar(label='Cluster')
plt.show()
```

The scatter plot shows clusters of data points based on the first two factors.

- Data points are grouped into three clusters, each representing distinct patterns within the accident data.
- Clusters indicate that certain factors, such as location and severity, may distinguish types of accidents, possibly related to urban vs. rural settings or severity levels.

1.5.2 Interpretation

Clustering reveals natural groupings, which could help identify accident types with similar profiles. These insights could be useful for targeted safety interventions based on the specific characteristics of each cluster.

1.6 Step 5: Discriminant Analysis

Using the clusters as the target variable, we apply Discriminant Analysis to develop a classification model. This model can help predict the cluster for new data points.

1.6.1 Confusion Matrix

The confusion matrix below shows the accuracy of the Discriminant Analysis model in predicting cluster membership.

```
[8]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
     from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import classification_report, ConfusionMatrixDisplay

# Preparing data
X = factors
y = clusters


# Splitting into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=0)

# Applying LDA
lda = LinearDiscriminantAnalysis()
lda.fit(X_train, y_train)
y_pred = lda.predict(X_test)

# Displaying confusion matrix
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.title('Discriminant Analysis - Confusion Matrix')
plt.show()

# Classification report
print(classification_report(y_test, y_pred))

```



	precision	recall	f1-score	support
0	1.00	1.00	1.00	300
1	1.00	1.00	1.00	227
2	1.00	1.00	1.00	315
accuracy			1.00	842
macro avg	1.00	1.00	1.00	842
weighted avg	1.00	1.00	1.00	842

The confusion matrix shows the model's accuracy in classifying data points into clusters.

- The model achieved 100% accuracy, with no misclassifications among clusters, indicating strong separation between the clusters.
- The classification report shows perfect precision, recall and F1 scores, reinforcing the model's reliability.

This high accuracy suggests that the clusters identified in Step 4 are well-defined, with distinct characteristics that the model can reliably predict. This classification model could potentially be applied to new data points to predict accident types based on the patterns observed.

1.7 Data Export for Power BI Visualisation

To enhance our data analysis in Power BI, we will now export the following:

- Original data with added PCA components (reduced dimensions): This enables pattern visualisation without all original variables.
- Cluster labels: Include labels assigned to each data point by the K-means model; this allows visualisation of each accident record's cluster category, aiding in profiling accident types.
- Factor scores (from factor analysis): Export the dataset with factor scores to visualise latent patterns identified by factor analysis.
- Classification labels (if using discriminant analysis): If discriminant analysis was used to classify clusters, add these predicted labels; this allows Power BI to categorise and filter new data points by predicted cluster.

These exports will provide Power BI with the necessary data to create insightful visualisations and analyses of accident patterns and classifications.

```
[9]: # Add PCA components (first 5 for example), clusters, and factors to the
      ↪ original data
data_export = data.copy()
data_export[['PC1', 'PC2', 'PC3', 'PC4', 'PC5']] = pca.transform(data_scaled)[:
      ↪, :5] # First 5 components
data_export['Cluster'] = clusters # Clustering labels
data_export[['Factor1', 'Factor2', 'Factor3', 'Factor4', 'Factor5']] = factors[:
      ↪, :5] # First 5 factor scores
```

```
# Generate classification labels based on factor scores
classification_labels = lda.predict(factors)

# Add classification labels to the DataFrame
data_export['Classification_Label'] = classification_labels

# Save to CSV for Power BI
data_export.to_csv('accident_analysis_export.csv', index=False)
```